

Capstone #2 -- Examining and Predicting Fake News

Introduction

The problem of fake online news is a persistent concern in contemporary society, impacting politics and society. Many NLP and machine learning scholars have researched this increasingly important problem and have developed highly-technical, insightful analyses. In this project, I examine the problem of fake news classification, by analyzing a large dataset of scraped news articles using various Python libraries. Using the large fake news dataset scraped by Maciej Szpakowski available at <https://github.com/several27/FakeNewsCorpus>, I sample a subset of news articles and perform text analysis on them. This project is relevant to people who are interested in how text analysis approaches can aid in the way we understand contemporary online journalism.

Decision makers in a wide array of professions rely on accurate journalistic information in their work. While the Internet enables access to a wealth of information, it is also a medium by which disinformation can be easily spread. Machine learning can provide a way to quickly identify sources as being suspect or reliable.

Data Cleaning and Pre-processing

Initial sampling and EDA

The corpus used includes over 20 million articles. Because of considerations regarding time and processing power, I first decided to sample 10,000 articles for my analysis. While performing initial EDA on the corpus, it became apparent that all articles were categorized into 12 different news types (Fig. 1).

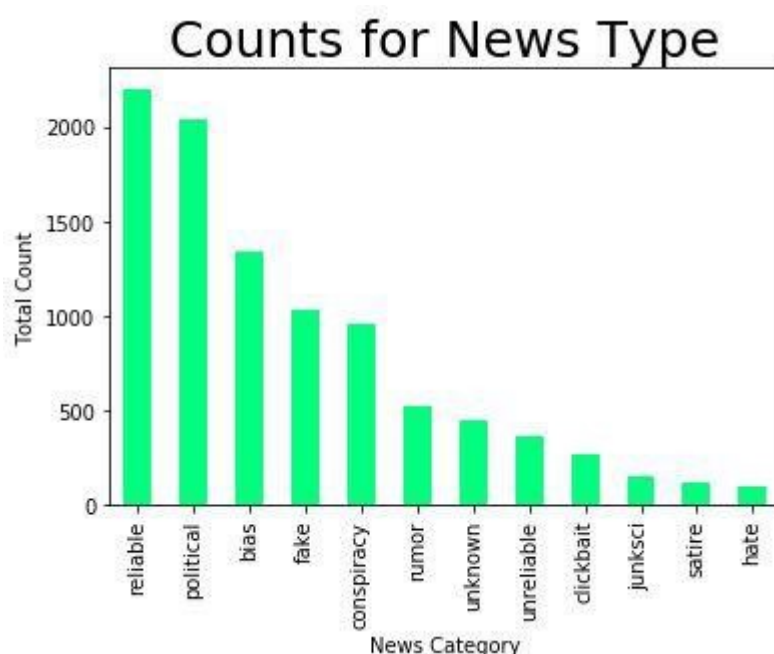



Fig 1.

Upon inspecting each of these news types, some issues began to emerge. Namely, all articles from any given source were given a particular label without consideration of individual articles. For example, all articles from *nytimes.com* were labelled as “reliable,” all articles from “beforeitsnews.com” were labelled as ‘fake’, and all articles from “sputniknews.com” were labelled as “bias”. Figure 2 shows the top ten news sources from particular web domains represented in the dataset.

Fig. 2



While the vast majority of articles from The New York Times can likely be considered reliable, some of the labelling of other various news sources seems to present some issues in the data. For example, all articles from ‘dailykos.com’ were labelled as ‘political’ (Fig. 5), all articles from ‘express.co.uk’ were labelled as ‘rumor’ (Fig. 6), and all articles from ‘sputniknews.com’ (Fig. 7) were labelled as ‘bias.’ Is every single article from each of these sources inherently more political or biased than articles from sources labelled as ‘reliable’? Because of the issues presented with the labelling, I decided to resample from the original dataset.

Fig. 3




Fig. 4

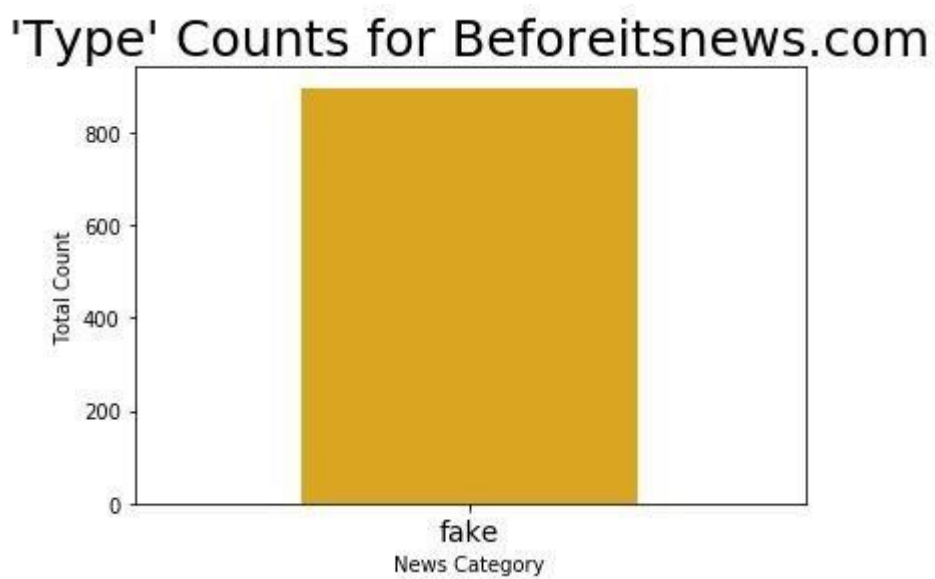


Fig. 5




Fig. 6

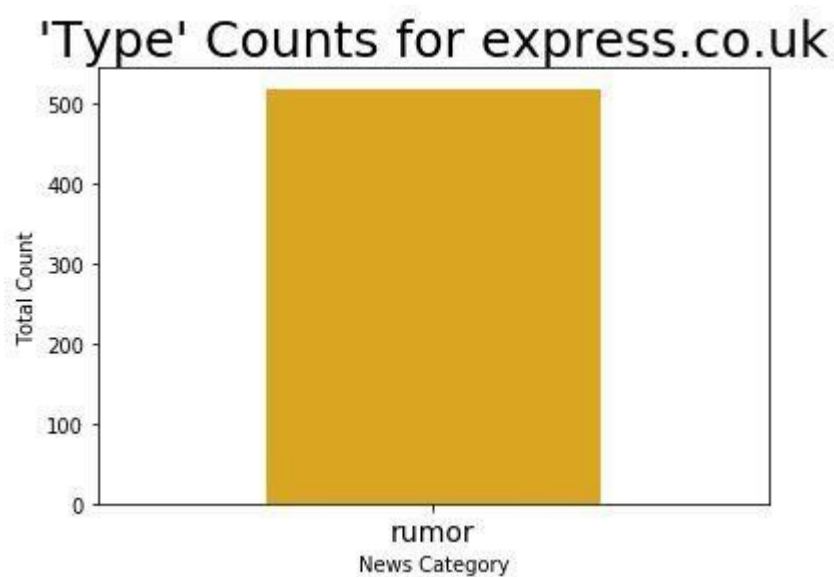



Fig. 7



Initial sampling

Due to the inherent biases in the dataset for the labelled categories 'junksci', 'bias', etc., it was initially decided to resample the dataset by including data from the two categories 'fake' and 'reliable'. Although there are issues with designating all articles from certain sources with the labels 'fake' or 'reliable', the sources used with these labels seem to be more consistent than with the other labels -- nytimes.com is, for the most part, reliable, whereas, a major website represented in the data, beforeitsnews.com, is not.

EDA of Sampled Data

The sampled data contains 10,000 articles labelled 'reliable' and 10,000 labelled 'fake'. The majority of these 20,000 articles were from nytimes.com (labelled as 'reliable') and beforeitsnews.com (labelled as 'fake'). There were, however, 141 difference online news sources represented in the data (Fig. 9).

Fig 8.





Fig. 9

News Source (domain)



Sentiment Analysis

The nltk Vader Sentiment Intensity Analyzer and TextBlob were used to analyze polarity in the corpora. Documents were labeled 1 (positive) if their score was greater than 0.2 (on the 1 to 1 scale used by both libraries), -1 (negative) if the score was less than -0.2, and neutral if it was between them.

Textblob was also used to calculate subjectivity scores. The scale for subjectivity scores is -1 to 1. Documents were labeled 1 (biased) if their subjectivity score was greater than 0.55, and -1 (unbiased) if the score was less than 0.45, and 0 if the score was between 0.45 and 0.55.

Polarity scores

	Label	SIA Polarity Score	TextBlob Polarity Score	TextBlob Subjectivity
Fake News	1	5468	983	1128
	-1	3013	67	5784
	0	5013	7944	3088
Reliable News	1	6721	1175	762
	-1	2581	73	6800
	0	698	8752	2438

The averages of all the scores taken for both the 'fake news' and 'reliable news' in the sampled data are as follows:

Fake

Textblob average polarity score: 0.105032

Textblob average subjectivity score: 0.4175212

SIA compound average: 0.22

SIA negative average: 0.072

SIA neutral average: 0.83

SIA positive average: 0.92

Reliable

Textblob average polarity score: 0.098

Textblob average subjectivity score: 0.407

SIA compound average: 0.377

SIA negative average: 0.059

SIA neutral average: 0.84

SIA positive average: 0.092

Additionally, 23% of the articles labeled 'fake' had a subjectivity score higher than or equal to .5, whereas only 17% of the articles labeled 'reliable' had such a value.

Predictive Modeling

The data was divided into a training and testing set. A several classifier was trained using scikit-learn, and then used to predict the labels for the testing data. Accuracy varied depending on vectorization approach:

Bag-of-words

Predictions were 86.7% accurate with Multinomial Naïve Bayes.

87% accurate with LinearSVC().

87.5% accurate with XGBoost().

Tf-idf

Predictions were 87.2% accurate with Multinomial Naïve Bayes.

90.8% accurate with LinearSVC().

89.0% accurate with XGBoost().

Tf-idf with two bigrams

Predictions were 90.4% accurate with Multinomial NB.

91.7% accurate with LinearSVC().

83.3% accurate with XGBClassifier().

There may be some issues with this model, linked with the data itself. Inspecting the most predictive features revealed, for example, that some of the most predictive bigrams for classification were 'york time' and 'york citi'. This is due to the over-representation of The New York Times as news labelled as 'reliable' in the dataset.

In order to deal with the problem of the over representation of articles from the New York Times in the dataset, the data was resampled to get a more varied distribution of article sources.

Analysis after Resampling

The New York Times and Beforeitsnew.com were vastly overrepresented in the dataset. While there were 141 different domains represented, the vast majority of documents come from one of two domains.

Fig. 10

'Domain' Counts in Initially Sampled Data





Fig 11.

'Domain' Counts in Undersampled Data



Because the domains were so unbalanced, undersampling was performed to somewhat rectify this problem. The New York Times and beforeitsnews.com were undersampled to 100 random articles. After this, the distribution of domains was more balanced.

Undersampling them left the dataset with only 3410 articles, which is a potential issue. Several classifiers were trained on the newly sampled data, and new accuracies were observed. Again, accuracy varied depending on the vectorization approach:

Bag-of-words

Predictions were 78.4% accurate with Multinomial NB.

80.7% accurate with LinearSVC().

79.8% accurate with XGBClassifier().

Tf-idf

Predictions were 63.9% accurate with Multinomial NB.

85.6% accurate with LinearSVC().

84.6% accurate with XGBClassifier().

Tf-idf with two bigrams

Predictions were 66.0% accurate with Multinomial NB.

82.1% accurate with LinearSVC().

78.7% accurate with XGBClassifier().

It is clear from looking at the informative features that document source is still playing a large role in predicting whether a given document is 'fake' or 'reliable'. Even after resampling, some of the most informative features for td-idf with bigrams were "new york", "associ press", and "thomson reuter". This indicates that document source is highly relevant in determining whether a document is predicting as being 'fake' or 'reliable'

Multi-class Classifier

Additionally, a multi-class classifier was built to include other classes than 'fake' and 'reliable' from the initial dataset. The categories of 'bias' and 'political' were also included, as these are the next largest types represented in the dataset. A multinomial Naïve Bayes, Linear SVC and XGBoost classifier was built, using `count_vectorizer()`, which yielded 68% accuracy. As the heat maps in Figures 12-14 shows, the classifier most accurately predicted the category of documents that had been labelled as "reliable".

Figure 12

Confusion matrix of MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)






Figure 13

Confusion matrix of LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True, intercept_scaling=1, loss='squared_hinge', max_iter=1000, multi_class='ovr', penalty='l2', random_state=None, tol=0.0001, verbose=0)

**Figure 14**

Confusion matrix of XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=3, min_child_weight=1, missing=None, n_estimators=100, n_jobs=1, nthread=None, objective='multi:softprob', random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None, silent=None, subsample=1, verbosity=1)



The same classifiers were then refit after using `tf_idf` vectorizer—the heatmaps of each are in figures 15-17:

Figure 15

Confusion matrix of `MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)`




Figure 16

Confusion matrix of `LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True, intercept_scaling=1, loss='squared_hinge', max_iter=1000, multi_class='ovr', penalty='l2', random_state=None, tol=0.0001, verbose=0)`





Figure 17

Confusion matrix of XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=3, min_child_weight=1, missing=None, n_estimators=100, n_jobs=1, nthread=None, objective='multi:softprob', random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None, silent=None, subsample=1, verbosity=1)



XGBClassifier() with tf-idf vectorization proved to be the most accurate with 75.7% accuracy.

All categories:

An additional analysis was done including all of the categories from the dataset, rather than the four most popular ones. The confusion matrices for each vectorization and classifier pair are shown in figures 18 to 23.

The results were less useful, due to the size of the sample of the dataset, as some categories had perfect precision and recall, due to their low representation in the dataset.

Conclusion

The most significant finding from the sentiment analysis of documents labelled 'fake' or 'reliable' from the dataset was that 23% of the articles labeled 'fake' had a subjectivity score higher than or equal to .5, whereas only 17% of the articles labeled 'reliable' had this high of a subjectivity score. It seemed that the TextBlob subjectivity score was more important than the polarity score in distinguishing between 'reliable' and 'fake' documents.

In the predictive analysis, various classifiers using different vectorization techniques were fairly accurate in determining whether a document from the testing data was labelled as 'fake' or 'reliable'. An inspection of the most informative features indicated, however, the terms linked to document source were particularly important in determining which class a given document belonged to. This indicates that words linked to what source an article came from may be more important than other words in determining whether it will be classified as 'fake' or 'reliable'.

Lastly, several classifiers were built to predict the four labels of 'fake', 'reliable', 'bias', and 'political', found in the dataset. The classifiers consistently performed best at predicting the 'reliable' label for documents in the testing data. The same classifiers were used to predict all of the labels in the sampled dataset. The results were less helpful because some of the categories were very underrepresented using previously made sample.

A next step for this project would be to consider using significantly more data. While this is time intensive in terms of processing, it may yield much more useful results for the multi-class classification.

An important takeaway of this project is that it is somewhat difficult to unlink fake news detection from human labeling. For example, even after the resampling and significant inclusion of other sources than the New York Times and beforeitsnews.com, some of the most predictive features in the most accurate classifier were "reuters", and "ap". Before the resampling, the most predictive bigrams were "new york", "main story", and other bigrams often repeated in New York Times articles. If all articles from a given source are labelled as "fake", "bias", or "junksci", a machine learning classifier may actually be more or less predicting the source, if certain phrases repeat in articles from a given source, than words and phrases that are inherently distinctive to fake or reliable news.